

STUDENT GUIDE

API and HPC User Guide

Configure WorkBuddy with the university API and start a GPU development session.

Official documentation

[WorkBuddy setup](#)

[API usage](#) | [Models and pricing](#)

[HPC platform login and connection](#)

[GPU development environments](#) | [GPU pricing](#)

Part 1 / API configuration in WorkBuddy

Before you begin

Connect to HKUST-GZ campus Wi-Fi or the campus VPN and have your assigned API key ready. Each student receives **RMB 50 in API credits** and **RMB 50 in GPU credits**.

1. Install and sign in

Download WorkBuddy from workbuddy.cn, install it, and sign in.

2. Add a custom model

Open the bottom-left profile menu, then **Settings** → **Models** → **Add Model**. Complete the four fields below. Click **Test connection**, then **Save**.

Field	Value
Provider	Custom
API Address / Base URL	https://aigc-api.hkust-gz.edu.cn/v1
API Key	Paste the API key assigned to you.
Model Name	DeepSeek-V4-Flash

Enter the model ID exactly, including capitalization.


Keep your key private. Do not include it in screenshots, shared files, or submitted code. Your API key and HPC login password are separate credentials.

PART 1 / WORKBUDDY SETUP

Configure your custom model


Use Settings to add the university API connection.

1. Open Settings → Models → Add Model




Models and Add Model in Settings.

2. Select Custom



Choose Custom in the provider list.

3. Enter your API details



4. Test and save

Use the base URL and model name from the configuration table on page 1. Paste your assigned API key into **API Key**.

Click **Test connection**. Once the connection succeeds, click **Save**.

To update an existing model, open its edit control. The same API fields appear in the Edit Model dialog.

Choose a model and test it


Use the model selector in WorkBuddy to activate your saved configuration.

3. Select your custom model

Open the **Tasks** view. Use the model selector at the lower-right of the chat box, scroll to the **Custom** group, and choose the model you added.

4. Send a short test

Send “Hi” and check that your selected model returns a reply.



Select your saved custom model.

API models and pricing

Model ID	Context	Input*	Output*
DeepSeek-V4-Flash	1M	0.75	2.25
DeepSeek-V4-Pro	64K	2.25	6.75
gpt-5.6-luna	1M	1.47	8.80
gpt-5.6-terra	1M	14.66	88.00
GLM-5.2	256K	4.00	14.00
Kimi-K3	128K	10.00	50.00

6 API models. Context indicates the maximum context-window size in tokens.

*RMB per 1 million tokens; prices rounded to two decimal places.

See the Models and pricing link at the beginning of this guide for the full catalog, current availability, and pricing.

PART 1 / OPTIONAL API ACCESS

Call the API from code

Use these alternatives when a script needs direct access to the same service.

Python / OpenAI SDK

Install the SDK with **`pip install openai`**. Replace **YOUR_API_KEY** in the example with the API key assigned to you, then run the code.

```
from openai import OpenAI

client = OpenAI(
    api_key="YOUR_API_KEY",
    base_url="https://aigc-api.hkust-gz.edu.cn/v1",
)
response = client.chat.completions.create(
    model="DeepSeek-V4-Flash",
    messages=[{"role": "user", "content": "Hello!"}],
)
print(response.choices[0].message.content)
```

Terminal / cURL (macOS or Linux shell)

Replace **YOUR_API_KEY** below with your assigned key before running.

```
curl -X POST "https://aigc-api.hkust-gz.edu.cn/v1/chat/completions" \
  -H "Authorization: Bearer YOUR_API_KEY" \
  -H "Content-Type: application/json" \
  -d '{"model": "DeepSeek-V4-Flash",
    "messages": [{"role": "user", "content": "Hello!"}]}'
```

Keep your key private.

Remove your actual key before sharing code or screenshots. Successful API requests consume API credits.


Sign in and open Model Develop

Access the university GPU development platform through your browser.

1. Open the teaching portal

<https://login.hpc.hkust-gz.edu.cn>

Use the HPC credentials assigned for your teaching access. The username field asks for the name without the @hkust-gz.edu.cn suffix. Select **ENG** if needed, then click **Login**.



The login panel features a header with a logo and the text "HKUST (GZ) HPC+AI Intelligence Computing Center Platform". In the top right corner, there are language options: "简体中文" and "ENG". The main form includes a "Username" field with the placeholder text "without(@hkust-gz.edu.cn)", a "Password" field, and a prominent "Login" button. At the bottom, a note states "Unified Identity Authentication for Campus Internal Use".

Figure 1. Login panel.




Figure 2. Model Develop icon.

2. Open Model Develop

On the platform desktop, double-click **Model Develop**, indicated by the red arrow. The Development Environment view opens.

Next step

Review the GPU prices on the next page, then create your VSCode environment on page 7.

PART 2 / GPU PRICING

GPU prices and billing

Rates are in RMB per card per hour for HPC I/II computing resources.

For VSCode development environments

The development center uses the low-tier rate and does not offer free debugging resources.

GPU / accelerator	Low priority	Medium priority	High priority
NVIDIA A800	1.60	2.50	4.00
NVIDIA A30	0.50	1.00	1.50
NVIDIA A40	0.50	1.00	1.50
Ascend	0.50	1.00	1.50

4 accelerator families in 2 price groups: A800 (1 family), and A30/A40/Ascend (3 families).

How usage is charged

Runtime is measured from actual start to end. Sessions shorter than one hour are charged by the second. For MIG-enabled GPUs, charges depend on the requested GPU memory.

GPU cost example

4 A800 cards × 1 hour × RMB 1.60 = **RMB 6.40**.

This is the GPU component only; separately billed resources are excluded.

Recommendation: make your credits last longer

Choose a lower-cost GPU, such as an A30 or A40, when it meets your task requirements and is available. The lower hourly rate lets your RMB 50 GPU credit cover more runtime. Request only the GPUs you need and stop the environment when finished.

Create a GPU environment

Choose a VSCode environment and the GPU resources available to your account.

3. Click Create Development Environment




Figure 3. Start a new environment.

4. Complete the form

Enter an environment name. Select **VSCode**, cluster **AI-X86_NVIDIA**, image **admin/webide-vscode:v6.2**, and **GPU** under Resource Specifications. Select the GPU option available for your course task.

A screenshot of the GPU configuration form. The form includes the following fields and options:

- Environment Name:** A text input field containing 'xxxx-kAIXDDL'. A red arrow points to this field.
- Environment Type:** A dropdown menu with 'VSCode' selected. A red arrow points to this dropdown.
- Select Cluster:** A dropdown menu with 'AI-X86_NVIDIA' selected. A red arrow points to this dropdown.
- Select Image:** A dropdown menu with 'admin/webide-vscode:v6.2' selected. A red arrow points to this dropdown. Above the dropdown is a filter 'VSCode' with a 'clear' link and a 'filterable' link.
- Resource Specifications:** Radio buttons for 'CPU' and 'GPU' (selected). Below them is a dropdown menu with 'ALL' selected. To the right of this dropdown is a text box showing '集群AI-X86_NVIDIA | 32核 | GPU卡4张(NVIDIA8...'. A red arrow points to this text box.
- Shared Memory (shm-size):** A text input field with the placeholder 'Please enter shared memory size in MB'.
- Custom Port:** A text input field with the placeholder 'Please enter port' and a '+' icon.
- Buttons:** 'Cancel' and 'Create Development Environment' buttons at the bottom right. A red arrow points to the 'Create Development Environment' button.

Figure 4. GPU configuration and the Create Development Environment button.

Choose resources deliberately. The configuration shown uses 4 GPUs and 32 CPU cores. Your allocation should match your course task. Use the course-specified resources, then click **Create Development Environment**.

Open VSCode and verify the GPU

Wait for the environment to run, then check the hardware inside the session.

5. Wait for RUN and open VSCode

When the environment shows **RUN**, click its **VSCode** tile. The card also provides a Web Terminal and a Stop control.




Figure 5. A running environment with the VSCode entry point.

6. Check GPU visibility

In VSCode, open a terminal and run:

```
nvidia-smi
```

PROBLEMS		OUTPUT		DEBUG CONSOLE		TERMINAL	
N/A	28C	P0		60W /	400W	0MiB / 81920MiB	0% Default Disabled
1	NVIDIA	A800-SXM4-80GB		On		00000000:25:00.0 Off	0
N/A	31C	P0		59W /	400W	0MiB / 81920MiB	0% Default Disabled
2	NVIDIA	A800-SXM4-80GB		On		00000000:50:00.0 Off	0
N/A	32C	P0		63W /	400W	0MiB / 81920MiB	0% Default Disabled
3	NVIDIA	A800-SXM4-80GB		On		00000000:55:00.0 Off	0
N/A	29C	P0		57W /	400W	0MiB / 81920MiB	0% Default Disabled
Processes:							
GPU	GI	CI		PID	Type	Process name	GPU Memory Usage
	ID	ID					
No running processes found							

Figure 6. NVIDIA GPU status in the VSCode terminal.

GPU rows indicate that NVIDIA devices are visible. “No running processes found” means no GPU process is listed at that moment. It does not indicate a failed allocation.

Finish the session cleanly

Save your work and stop the environment when you finish.

7. Save your work and stop the environment

Save your files and keep a backup of important results. Return to **Model Develop** and click **Stop** on the environment card when you finish. Check that it no longer shows RUN; do not rely on closing the browser tab to stop it.

Troubleshooting: 4 common issues

Symptom	What to check
WorkBuddy cannot connect	Check campus Wi-Fi/VPN, the base URL, and your assigned API key.
Model is missing or rejected	Save the model first, select it under Custom, and check its exact ID.
Environment does not reach RUN	Inspect Environment Log and check the resource request before retrying.
GPU check fails	Confirm that you selected GPU and ran the command inside the running environment.